

CreatorPulse PRD

■ Abstract

CreatorPulse is an AI-powered assistant that automates daily newsletter research and drafting for creators and agencies. It aggregates trusted sources, surfaces emerging trends, and generates a voice-matched draft newsletter—delivered directly via email each morning. The goal is to reduce newsletter creation time from hours to minutes while improving content consistency and engagement.

■ Business Objectives

- Cut newsletter drafting time from 2–3 hours to under 20 minutes.
- Achieve at least a 70% draft-acceptance rate within 90 days.
- Increase engagement (open and click rates) for 60% of active users.

■ KPI

GOAL	METRIC	QUESTION
Reduce drafting time	Avg. review time per accepted draft ≤ 20 min	Are users saving time?
Increase adoption	Draft acceptance rate $\geq 70\%$	Do users trust the AI drafts?
Improve engagement	Open/CTR $\geq 2\times$ baseline	Is the newsletter performing better?

■ Success Criteria

A successful MVP will have at least 100 active users achieving >70% draft acceptance and measurable engagement uplift within 12 weeks.

■■■■ User Journeys

A content curator connects their Twitter, YouTube, and newsletter feeds. Each morning, they receive an AI-generated newsletter draft with curated insights and emerging trends. They review, tweak, and send the newsletter within 20 minutes—all via email.

■ Functional Requirements

- Source Connections: Integrate with Twitter, YouTube, and RSS feeds.
- Trend Engine: Detect emerging topics and highlight spikes in relevance.
- Writing Style Trainer: In-context learning from uploaded newsletters (private per user).
- Newsletter Draft Generator: Create full draft with intro, curated links, commentary, and trend highlights.

- Morning Delivery: Deliver draft and trends digest at 8:00 local via email.
- Feedback Loop: 🟡 / 🟢 reactions and edit tracking to improve output quality.

■ **Model Requirements**

SPECIFICATION	REQUIREMENT	RATIONALE
Model Type	Prompt-based LLM (GPT-class)	Fast adaptation, no fine-tuning cost.
Latency	P95 < 5s per draft	Smooth email delivery experience.
Context Window	8K tokens	Support multiple source summaries.
Fine Tuning	Not required for v1	In-context training suffices.

■■ **Risks & Mitigations**

RISK	MITIGATION
API rate limits	Implement caching and delta crawls.
Voice mismatch	Human feedback loop and retraining pipeline.
False trends	Use ensemble detection and manual override flags.
Email deliverability	Verified sender domains and batch sends.

■ **Future Expansion (v2)**

- Optional responsive web dashboard for managing sources and analytics.
- Deeper integrations with Google Trends, arXiv, and niche industry blogs.
- Multi-language draft generation.
- Social publishing to X and LinkedIn.